

GPGPU - Project report

Nicolas Portal - Alexandre Yvart - Geoffrey Jount

14 juin 2020

Table des matières

1	Introduction	1
2	Problem description	1
3	CPU implementation	1
3.1	Issue 1	1
3.2	Issue 2	2
3.3	Issue 3	2
3.4	Issue 4	2
4	optimization	3
5	Output segmented images	3
6	GPU implementation	4
7	Bottlenecks	5
8	Enhancements	6
8.1	Enhancement 1	6
8.2	Enhancement 2	7
9	Benchmarks recapitulation	7
10	Tasks repartition	7

1 Introduction

For educational purpose in the GPGPU subject, we are asked to work on a project using a GPU parallel method. We will use the CUDA library.

2 Problem description

This project's goal is to implement an image segmentation algorithm : the max-flow/min-cut algorithm. A picture containing numerous pixels however, processing one after another in a simple `for` loop would show to be very slow. We thus need to make the process parallel, and implement it on GPU for an even faster result : the *push-relabel* algorithm, parallel version of the first mentioned, is what we seek.

3 CPU implementation

We followed the guidelines found on the NVIDIA presentation.
All in all there are eight matrices :

- for the heights of each node
- for the excess flow of each node
- for the links between the source and each node
- for the links between the sink and each nodes

And the four last matrices for the links between each nodes : top, bottom, right and left neighbours matrices were labeled with the inverse of the euclidean distance in RGB space.

The links between the source and each node and the links between the sink and each node were labelled with the inverse of the average distance to the marked pixels in RGB space. Each node were linked to the source and sink. The height of each node was 0 except for the source which had a height equals to the number of nodes. The `std_load` library was used to load `jpg` images both in grayscale and RGB color space. Images were stored as one dimensional vectors with contiguous red, green and blue pixels. the size of these vectors were three times as large as the number of pixels in an image. The `ppm` file format was used to generate the output segmented image.

3.1 Issue 1

Our implementation soon ran into an infinite loop which remained one of the core issue for several weeks. As a result, the code was changed to make it possible for each node to push towards the sink or the source. This did not solved the issue. The weights of links between the source and sink on one hand and every node on the other hand was changed using the histogram on red, green and blue channel of the marked pixels to compute a probability. This did not solved the issue as well.

It was found that there was actually no infinite loop but the algorithm was very long to terminate as two nodes would often exchange flow between them waiting for one of them to reach the maximum height. To solve this glitch, it was decided to manually lower the maximum height. The maximum height was no longer set to the number of nodes but rather to a low number such as five or ten. This made it possible for our algorithm to terminate.

maxHeight	time
5	146ms
10	271ms
100	1755ms

3.2 Issue 2

An other issue quickly appeared. Using a depth first search to find nodes yielded bad results as every pixel would be set to white.

When it comes to the way to get the output image, it was decided to set every pixel with a height above zero to one. This did not generate optimal results but it was good enough to make out the segmentation.

A small eight by eight image was used to make the debug step simpler. This image was a ppm image with every pixels on the top part of the image set to blue and every pixel on the lower part of the image set to red. One pixel was set to black and white in each region.

3.3 Issue 3

Although the maximum height was decrease to five or ten, the run time was still to high. It took 15 to 20 minutes for the algorithm to terminate.

Lowering the maximum height resulted in a 15, 20 minute running time. This was still too high. It was found out that calling the push and relabel function on every nodes would decrease the run time significantly. The check for active nodes was now done by the push and relabel function. It was no longer needed to write a separate function to retrieve the right active nodes. It was also no longer required to loop on every pixels to find active nodes. This decreased the run time to 10 seconds which was a huge improvement.

3.4 Issue 4

For a long time we simply checked the height value of each pixel to get the min cut of the graph and thus decide the color of each pixel.

Since we were not satisfied with our output segmented image, it was decided

to change the way we recovered the min cut. Although there were very little explanation about the way the BFS algorithm should have been implemented in NVIDIA slides, the algorithm was successfully implemented. It was a little bit of a disappointment though since this did not change the look of the output segmented image. This led us to understand that this non optimal segmented image resulted from a flaw in the core algorithm (maybe the weights of the links).

4 optimization


It was found out that the count active function slowed down the algorithm.

In order to make it quicker, It was decided to no longer loop on the whole excess flow matrix but rather, to return true as soon as the function found an active node. This slight change required to change the return type of the function from int to bool. It was not necessary to count every active nodes anymore. The program would keep in the main while loop as long as the count active function returned true. This resulted in a slight but noticeable decrease in run time.

maxHeight = 5	<u>Before</u>	<u>After</u>
Time	159 ms	146 ms

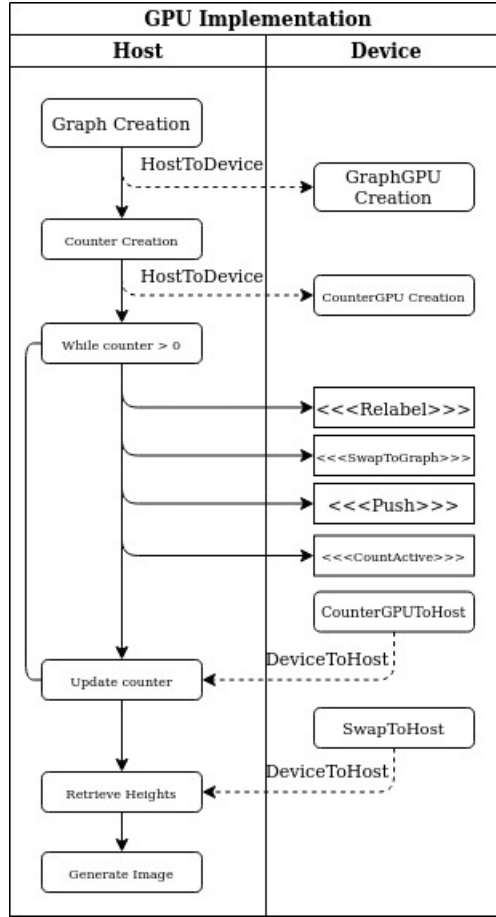
5 Output segmented images





6 GPU implementation

The first GPU implementation is kept as close as possible to the CPU version. The graph structure is instantiated on the CPU just like the CPU version, except that we flattened the matrices on one dimension. The graph structure is then copied on the device, with the addition of a swap matrix used for the relabeling step. Then, the push and relabel kernels are called one after the other, where each thread is representing one node of the graph. The final cut is done just like the CPU version. We didn't have much issues on implementing the GPU version since the CPU version was already functional.



7 Bottlenecks

After proceeding to some profiling, we quickly saw that our kernel used to count_active nodes was a bottleneck because it was taking too much time despite not doing any calculus. Otherwise, most of the bottlenecks are on the allocation of memory on the device which is hard to optimize unless by allocation the whole structure at once. The push kernel is also a bottleneck since it's checking a lot of conditions. Overall, except from reducing the number of calls to cudaMalloc, we can only solve the bottlenecks with algorithmic optimizations.

```
==98771== Profiling result:
```

Type	Time(%)	Time	Calls	Avg	Min	Max	Name
GPU activities:	32.93%	2.8255ms	70	40.363us	24.543us	97.596us	push(GraphGPU*)
	22.75%	1.9516ms	71	27.487us	26.718us	30.046us	count_active(GraphGPU*, int*)
	22.18%	1.9030ms	70	27.185us	19.167us	48.478us	relabel(GraphGPU*, int*)
	13.15%	1.1279ms	70	16.113us	15.552us	16.671us	swap_to_graph(GraphGPU*, int*)
	5.70%	489.20us	19	25.747us	1.1840us	53.310us	[CUDA memcpy HtoD]
	1.80%	154.62us	72	2.1470us	1.4400us	49.982us	[CUDA memcpy DtoH]
	1.10%	94.526us	70	1.3500us	1.3440us	1.3760us	[CUDA memset]
	0.21%	18.111us	1	18.111us	18.111us	18.111us	graph_to_swap(GraphGPU*, int*)
	0.18%	15.583us	1	15.583us	15.583us	15.583us	setImage(GraphGPU*, int*)
	94.21%	218.61ms	12	18.217ms	5.1710us	218.29ms	cudaMalloc
API calls:	4.00%	9.2825ms	292	31.789us	1.2650us	102.25us	cudaDeviceSynchronize
	0.97%	2.2566ms	91	24.797us	3.5060us	142.50us	cudaMemcpy
	0.49%	1.1474ms	283	4.0540us	3.2650us	23.019us	cudaLaunchKernel
	0.09%	201.00us	70	2.8710us	2.6600us	6.0160us	cudaMemset

8 Enhancements

8.1 Enhancement 1

We reduced the time spent on the `count_active` kernel by removing the atomic operations. We didn't need to count how many nodes were active but instead we just had to check if none were active. So, by using a boolean we allowed concurrent accesses to the variable and reduced the time spent on counting nodes by half.

```
==104895== Profiling result:
```

Type	Time(%)	Time	Calls	Avg	Min	Max	Name
GPU activities:	39.36%	3.1715ms	70	45.306us	27.358us	98.620us	push(GraphGPU*)
	23.99%	1.9334ms	70	27.619us	19.551us	48.413us	relabel(GraphGPU*, int*)
	14.08%	1.1347ms	70	16.209us	15.839us	16.735us	swap_to_graph(GraphGPU*, int*)
	13.04%	1.0507ms	71	14.799us	12.831us	31.199us	count_active(GraphGPU*, int*)
	6.06%	488.56us	19	25.713us	544ns	53.278us	[CUDA memcpy HtoD]
	1.92%	154.94us	72	2.1510us	1.4710us	49.950us	[CUDA memcpy DtoH]
	1.05%	84.860us	70	1.2120us	1.0560us	1.3760us	[CUDA memset]
	0.25%	19.775us	1	19.775us	19.775us	19.775us	setImage(GraphGPU*, int*)
	0.24%	19.168us	1	19.168us	19.168us	19.168us	graph_to_swap(GraphGPU*, int*)
	91.37%	139.24ms	12	11.604ms	5.0940us	138.93ms	cudaMalloc
API calls:	5.76%	8.7764ms	292	30.056us	1.2010us	102.67us	cudaDeviceSynchronize
	1.48%	2.2536ms	91	24.764us	3.2820us	138.10us	cudaMemcpy
	0.91%	1.3831ms	283	4.8870us	3.0460us	42.539us	cudaLaunchKernel
	0.13%	201.96us	70	2.8850us	2.6810us	5.0890us	cudaMemset

8.2 Enhancement 2

Since most of the time spent by the API calls is on cudaMalloc, we reduced the number of allocation by 1 to gain 10 ms of execution time.

==108136== Profiling result:								
	Type	Time(%)	Time	Calls	Avg	Min	Max	Name
GPU activities:		37.51%	2.8582ms	70	40.831us	28.895us	88.124us	push(GraphGPU*)
		26.18%	1.9945ms	70	28.492us	24.063us	44.734us	relabel(GraphGPU*, int*)
		14.18%	1.0806ms	70	15.436us	14.975us	16.223us	swap_to_graph(GraphGPU*, int*)
		12.04%	917.59us	71	12.923us	10.975us	30.495us	count_active(GraphGPU*, int*)
		6.45%	491.05us	19	25.844us	1.1840us	53.726us	[CUDA memcpy HtoD]
		1.96%	149.34us	72	2.0740us	1.1840us	49.950us	[CUDA memcpy DtoH]
		1.18%	90.045us	70	1.2860us	928ns	1.3120us	[CUDA memset]
		0.25%	19.327us	1	19.327us	19.327us	19.327us	setImage(GraphGPU*, int*)
		0.24%	18.207us	1	18.207us	18.207us	18.207us	graph_to_swap(GraphGPU*, int*)
API calls:		79.06%	119.78ms	11	10.889ms	4.9440us	119.55ms	cudaMalloc
		17.77%	26.927ms	292	92.214us	1.2590us	142.54us	cudaDeviceSynchronize
		1.59%	2.4044ms	91	26.422us	3.5270us	144.97us	cudaMemcpy
		1.13%	1.7124ms	283	6.0500us	3.2100us	18.558us	cudaLaunchKernel
		0.15%	221.33us	70	3.1610us	2.7270us	15.903us	cudaMemset

9 Benchmarks recapitulation

Benchmark	Time	CPU Iterations	UserCounters...
BM_cpu/real_time	264 ms	264 ms	3 frame_rate=3.78102/s
BM_gpu/real_time	38 ms	38 ms	19 frame_rate=26.1102/s

10 Tasks repartition

Nicolas

- CPU implementation and debug
- Report
- GPU debug

Alexandre

- CPU debug
- GPU implementation

Geoffrey

- CPU debug
- Benchmarking implementation